

2016-06-14 Exam Exercises

The following exercises and their solutions were originally authored by *Professor Maria de Marsico*, for the "Basi di Dati" exam session on June, the 14th, 2016.

Their translation is ongoing.

1. Relational Algebra

Consider a database with the following schema, describing a manufacturer's operations:

```
PRODUCTS (ID, Description, UnitPrice)
WAREHOUSES (ID, Address)
STOCKS (ProductID, WarehouseID, Units)
```



STOCKS instances describe how many *products' units* are stocked and in which *warehouses*. When a product is stocked, its **Units** ≥ 1 .

Write **relational algebra expressions** for the following queries:

- For each product whose stock is equal or larger than 10 units, *in any warehouse*, get both:
 - the product's data (**ID**, **Description** and **UnitPrice**)
 - the addresses of all the warehouses where at least 10 product units are stocked.
- Find the **ID**, **Description** and **UnitPrice** of the products that *aren't* stocked at all, anywhere.



You can load a sample Relax dataset with this gist ID:
126fcdb8c1bedc5080270dff5f642186

1.1. Answer

We must identify the relations containing the needed data. We require all of them, as:

- **STOCKS** contains the number of stocked units for each product
- **PRODUCTS** holds the data for products details
- **WAREHOUSES** includes the addresses

The easiest query involves:

1. joining the three relations together, via **theta joins** where appropriate
2. performing a **selection** on the resulting relation, by filtering those tuples whose **Units** are equal or higher than 10
3. using a **projection** to pick out the attribute values we require



PRODUCTS and **WAREHOUSES** both feature an **ID** attribute, although these identify tuples in different relations, with different meanings. It wouldn't make sense to perform a **natural join** between them.

Let r identify the desired data:

$$r = \sigma_{Units \geq 10} \text{STOCKS} \bowtie_{WarehouseID=ID} \text{WAREHOUSES} \bowtie_{ProductID=PRODUCTS.ID} \text{PRODUCTS}$$

A less efficient *alternative*, due to more *joins*, could be:

$$r = \sigma_{Units \geq 10} (\text{WAREHOUSES} \bowtie_{ID=WarehouseID} \text{STOCKS} \bowtie_{ProductID=PRODUCTS.ID} \text{PRODUCTS})$$

We then need to select the relevant attributes, via a **projection** on r :

$$\pi_{PRODUCTS.ID, Description, UnitPrice, Address}(r)$$

```

r = σ Units ≥ 10 STOCKS ⋈ WarehouseID = ID WAREHOUSES ⋈ ProductID =
PRODUCTS.ID PRODUCTS
π PRODUCTS.ID, Description, UnitPrice, Address (r)

```

1.2. Answer



This class of relational algebra problems is best handled with **subtractions**. The tuples that *don't* meet the selection criteria are first collected and then removed from the set of all the candidate tuples.

In this case we don't need to query the **WAREHOUSES** relation, seeing as it contains no relevant data for our purposes.



Products that *aren't* stocked *don't appear* in **STOCKS** instances; there are no such tuples whose **Units** value is **0**.

Let **r** be the relation which includes the data of all those products we aren't interested in:

$$r = \pi_{ID, Description, UnitPrice} (PRODUCTS \bowtie_{ID=ProductID} STOCKS)$$

We are selecting **all** the tuples that match stocked products, referenced in **STOCKS** via the **ProductID** attribute. *Unstocked* products, absent from **STOCKS**, won't be included in the *join*.

We finally **subtract** the data of all stocked products, **r**, from the set of all products (stocked and otherwise):

$$PRODUCTS - r$$



The initial **projection** ensures that the two relations' schemas are **compatible**, as required by the **subtraction**.

```
PRODUCTS -  $\pi$  ID, Description, UnitPrice (PRODUCTS  $\bowtie$  ID = ProductID STOCKS)
```